

Assignment 6.2

July 9, 2021

1 Assignment 6.2

1.1 Assignment 6.2.a

Using section 5.2 in Deep Learning with Python as a guide, create a ConvNet model that classifies images CIFAR10 small images classification dataset. Do not use dropout or data-augmentation in this part. Save the model, predictions, metrics, and validation plots in the dsc650/assignments/assignment06/results directory. If you are using JupyterHub, you can include those plots in your Jupyter notebook.

```
[1]: import keras
import tensorflow as tf
import os
from keras.datasets import cifar10
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd

from keras.preprocessing.image import ImageDataGenerator

import sklearn.metrics as metrics
from seaborn import heatmap
```

```
[2]: (X_train, y_train), (X_test, y_test) = cifar10.load_data()
```


1.2 Data Exploration

```
[3]: axs = [f'ax{i}' for i in range(0,10)]

fig, [axs[0:5],
      axs[5:10]] = plt.subplots(2,5, figsize=(16,9))

for ax in axs:
    ax.imshow(X_train[np.random.randint(0,10000)])

plt.show()
```



```
[4]: print('X_train shape:', X_train.shape)
      print('y_train shape:', y_train.shape)

      print(X_train.shape[0], 'train samples')
      print(X_test.shape[0], 'test samples')
```

```
X_train shape: (50000, 32, 32, 3)
y_train shape: (50000, 1)
50000 train samples
10000 test samples
```

1.3 Data Preprocessing

```
[5]: X_train = X_train.astype('float32')
      X_train /= 255

      X_test = X_test.astype('float32')
      X_test /= 255

      X_val_train = X_train[:10000]
      X_train = X_train[10000:]

      # Convert target data to single array of shape (50000,) and (10000,)
      y_train = y_train.reshape(y_train.shape[0])
      y_test = y_test.reshape(y_test.shape[0])

      y_val_train = y_train[:10000]
      y_train = y_train[10000:]
```

1.4 ConvNet Model

```
[6]: batch_size = 128
     epochs = 200
```

```
[7]: from keras.models import Sequential
     from keras.layers import Conv2D, Activation, MaxPooling2D, Dropout, Dense, Flatten
```

```
[8]: model = Sequential()
     model.add(Conv2D(32, (3, 3),
                     padding='same',
                     input_shape=X_train.shape[1:]))
     model.add(MaxPooling2D((2, 2)))
     model.add(Conv2D(64, (3, 3), activation='relu'))
     model.add(MaxPooling2D((2, 2)))
     model.add(Conv2D(128, (3, 3), activation='relu'))
     model.add(MaxPooling2D((2, 2)))
     model.add(Flatten())
     model.add(Dense(512, activation='relu'))
     model.add(Dense(10, activation='sigmoid'))

     model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 32, 32, 32)	896
max_pooling2d (MaxPooling2D)	(None, 16, 16, 32)	0
conv2d_1 (Conv2D)	(None, 14, 14, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 7, 7, 64)	0
conv2d_2 (Conv2D)	(None, 5, 5, 128)	73856
max_pooling2d_2 (MaxPooling2D)	(None, 2, 2, 128)	0
flatten (Flatten)	(None, 512)	0
dense (Dense)	(None, 512)	262656
dense_1 (Dense)	(None, 10)	5130

Total params: 361,034
Trainable params: 361,034

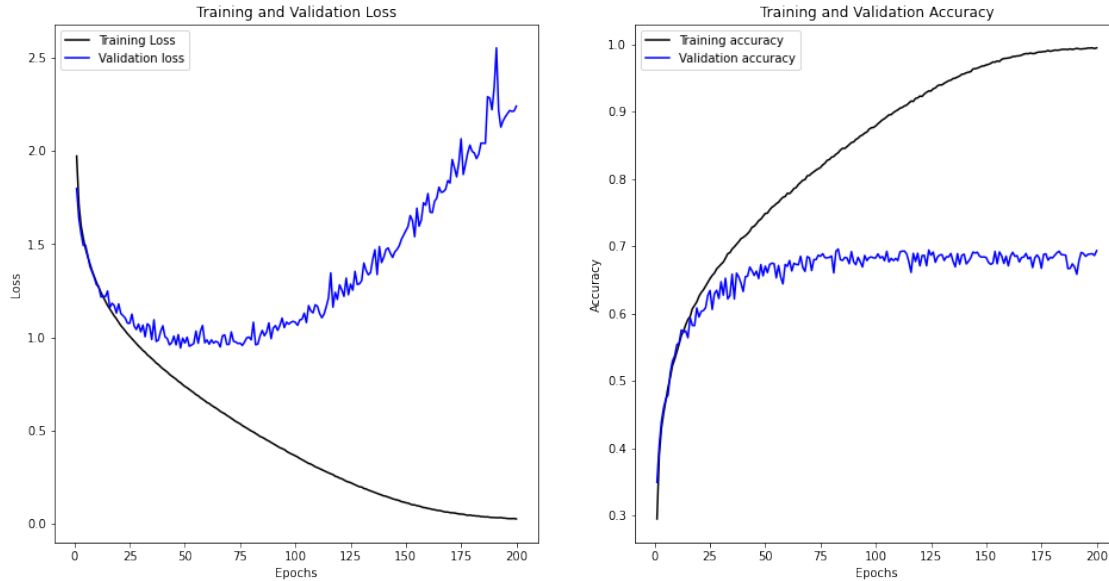
Non-trainable params: 0

```
[9]: model.compile(loss='sparse_categorical_crossentropy',  
                  optimizer=keras.optimizers.RMSprop(learning_rate=0.0001,  
↳ decay=1e-6),  
                  metrics=['acc'])
```

```
[10]: history = model.fit(X_train,  
                          y_train,  
                          batch_size=batch_size,  
                          epochs = epochs,  
                          validation_data=(X_val_train, y_val_train),  
                          verbose=False)
```

```
[11]: results = model.evaluate(X_test, y_test)  
print(results)  
  
history_dict = history.history  
  
acc = history_dict['acc']  
val_acc = history_dict['val_acc']  
loss_values = history_dict['loss']  
val_loss_values = history_dict['val_loss']  
epochs = range(1, len(acc) + 1)  
  
# Plotting metrics  
fig, [ax1, ax2] = plt.subplots(1, 2, figsize=(16, 8))  
  
ax1.plot(epochs, loss_values, 'k', label = 'Training Loss')  
ax1.plot(epochs, val_loss_values, 'b', label = 'Validation loss')  
ax1.set_title('Training and Validation Loss')  
ax1.set_xlabel("Epochs")  
ax1.set_ylabel("Loss")  
ax1.legend()  
  
ax2.plot(epochs, acc, 'k', label = 'Training accuracy')  
ax2.plot(epochs, val_acc, 'b', label = 'Validation accuracy')  
ax2.set_title('Training and Validation Accuracy')  
ax2.set_xlabel("Epochs")  
ax2.set_ylabel("Accuracy")  
ax2.legend()  
plt.show()
```

```
313/313 [=====] - 2s 5ms/step - loss: 2.3716 - acc:  
0.6730  
[2.371561050415039, 0.6729999780654907]
```




```
[12]: batch_size = 128
epochs = 60

model = Sequential()
model.add(Conv2D(32, (3, 3),
                padding='same',
                input_shape=X_train.shape[1:]))
model.add(MaxPooling2D((2, 2)))
model.add(Conv2D(64, (3, 3), activation='relu'))
model.add(MaxPooling2D((2, 2)))
model.add(Conv2D(128, (3, 3), activation='relu'))
model.add(MaxPooling2D((2, 2)))
model.add(Flatten())
model.add(Dense(512, activation='relu'))
model.add(Dense(10, activation='sigmoid'))

model.compile(loss='sparse_categorical_crossentropy',
              optimizer=keras.optimizers.RMSprop(learning_rate=0.0001,
              ↪decay=1e-6),
              metrics=['acc'])

history = model.fit(X_train,
                    y_train,
                    batch_size=batch_size,
                    epochs = epochs,
                    validation_data=(X_val_train, y_val_train),
                    verbose=False)
```

```
[13]: predictions = model.predict(X_test)
confusion_matrix = metrics.confusion_matrix(y_true=y_test, y_pred=predictions.
→argmax(axis=1))
fig, ax = plt.subplots(figsize=(16,10))
heatmap(confusion_matrix, annot = True, cmap='Blues', fmt='g', ax = ax);
```



```
[14]: predictions.argmax(axis = 1).shape
```

```
[14]: (10000,)
```

```
[15]: labels = ['Airplane', 'Automobile', 'Bird', 'Cat', 'Deer', 'Dog', 'Frog', 'Horse', 'Ship', 'Truck']
```

```
df = pd.DataFrame()
df['prediction'] = predictions.argmax(axis=1)
df['actual'] = y_test
df['prediction'] = df['prediction'].apply(lambda x: labels[x])
df['actual'] = df['actual'].apply(lambda x: labels[x])
df.to_csv("results/model_6_a_predictions.csv")
df.sample(20)
```

```
[15]:
```

	prediction	actual
1384	Truck	Truck
1264	Ship	Ship
7653	Ship	Ship
2146	Bird	Bird
4286	Bird	Bird
3676	Truck	Truck
8386	Cat	Cat
7997	Truck	Automobile
2335	Deer	Dog
2559	Horse	Horse
3497	Airplane	Airplane
8016	Deer	Deer
1939	Horse	Frog
935	Bird	Bird
465	Deer	Deer
3646	Cat	Cat
3160	Frog	Frog
2767	Horse	Dog
9019	Truck	Truck
3047	Cat	Frog

```
[16]: model.save('results/model_6_2_a.h5')
```

1.5 Assignment 6.2.b

Using section 5.2 in Deep Learning with Python as a guide, create a ConvNet model that classifies images CIFAR10 small images classification dataset. This time includes dropout and data-augmentation. Save the model, predictions, metrics, and validation plots in the dsc650/assignments/assignment06/results directory. If you are using JupyterHub, you can include those plots in your Jupyter notebook.

```
[17]: model = Sequential()
model.add(Conv2D(32, (3, 3), padding='same',
                input_shape=X_train.shape[1:]))
model.add(Activation('relu'))
model.add(Conv2D(32, (3, 3)))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))
model.add(Conv2D(64, (3, 3), padding='same'))
model.add(Activation('relu'))
model.add(Conv2D(64, (3, 3)))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))
model.add(Flatten())
model.add(Dense(512))
```

```

model.add(Activation('relu'))
model.add(Dropout(0.5))
model.add(Dense(10))
model.add(Activation('softmax'))

model.summary()

```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
conv2d_6 (Conv2D)	(None, 32, 32, 32)	896
activation (Activation)	(None, 32, 32, 32)	0
conv2d_7 (Conv2D)	(None, 30, 30, 32)	9248
activation_1 (Activation)	(None, 30, 30, 32)	0
max_pooling2d_6 (MaxPooling2D)	(None, 15, 15, 32)	0
dropout (Dropout)	(None, 15, 15, 32)	0
conv2d_8 (Conv2D)	(None, 15, 15, 64)	18496
activation_2 (Activation)	(None, 15, 15, 64)	0
conv2d_9 (Conv2D)	(None, 13, 13, 64)	36928
activation_3 (Activation)	(None, 13, 13, 64)	0
max_pooling2d_7 (MaxPooling2D)	(None, 6, 6, 64)	0
dropout_1 (Dropout)	(None, 6, 6, 64)	0
flatten_2 (Flatten)	(None, 2304)	0
dense_4 (Dense)	(None, 512)	1180160
activation_4 (Activation)	(None, 512)	0
dropout_2 (Dropout)	(None, 512)	0
dense_5 (Dense)	(None, 10)	5130
activation_5 (Activation)	(None, 10)	0
Total params: 1,250,858		

Trainable params: 1,250,858

Non-trainable params: 0

```
[18]: model.compile(loss='sparse_categorical_crossentropy',  
                  optimizer=keras.optimizers.RMSprop(learning_rate=0.0001,  
                  ↳decay=1e-6),  
                  metrics=['acc'])
```

```
[19]: batch_size = 128  
      epochs = 300
```

```
[20]: datagen = ImageDataGenerator(  
        rotation_range=40,  
        width_shift_range=0.2,  
        height_shift_range=0.2,  
        shear_range=0.2,  
        zoom_range=0.2,  
        horizontal_flip=True)  
  
datagen.fit(X_train)  
  
history = model.fit(datagen.flow(X_train,  
                                y_train,  
                                batch_size=batch_size),  
                    epochs=epochs,  
                    validation_data=(X_val_train, y_val_train),  
                    workers=4,  
                    verbose=False)
```

```
[21]: results = model.evaluate(X_test, y_test)  
      print(results)  
  
history_dict = history.history  
  
acc = history_dict['acc']  
val_acc = history_dict['val_acc']  
loss_values = history_dict['loss']  
val_loss_values = history_dict['val_loss']  
epochs = range(1, len(acc) + 1)  
  
# Plotting metrics  
fig, [ax1, ax2] = plt.subplots(1, 2, figsize=(16, 8))  
  
ax1.plot(epochs, loss_values, 'k', label = 'Training Loss')  
ax1.plot(epochs, val_loss_values, 'b', label = 'Validation loss')  
ax1.set_title('Training and Validation Loss')
```

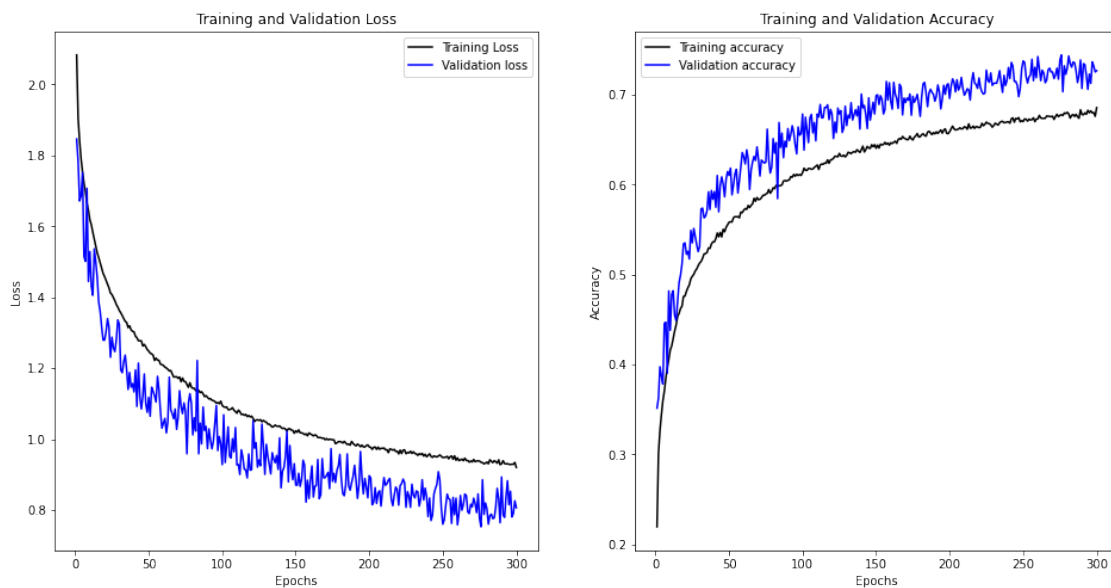
```

ax1.set_xlabel("Epochs")
ax1.set_ylabel("Loss")
ax1.legend()

ax2.plot(epochs, acc, 'k', label = 'Training accuracy')
ax2.plot(epochs, val_acc, 'b', label = 'Validation accuracy')
ax2.set_title('Training and Validation Accuracy')
ax2.set_xlabel("Epochs")
ax2.set_ylabel("Accuracy")
ax2.legend()
plt.show()

```


313/313 [=====] - 2s 8ms/step - loss: 0.8179 - acc: 0.7187
 [0.8179059028625488, 0.7186999917030334]



```

[22]: predictions = model.predict(X_test)
confusion_matrix = metrics.confusion_matrix(y_true=y_test, y_pred=predictions.
    ↪argmax(axis=1))
fig, ax = plt.subplots(figsize=(16,10))
heatmap(confusion_matrix, annot = True, cmap='Blues', fmt='g', ax = ax);

```



```
[23]: labels = ['Airplane', 'Automobile', 'Bird', 'Cat', 'Deer', 'Dog', 'Frog', 'Horse', 'Ship', 'Truck']
```

```
df = pd.DataFrame()
df['prediction'] = predictions.argmax(axis=1)
df['actual'] = y_test
df['prediction'] = df['prediction'].apply(lambda x: labels[x])
df['actual'] = df['actual'].apply(lambda x: labels[x])
df.to_csv("results/model_6_b_predictions.csv")
df.sample(20)
```

```
[23]:
```

	prediction	actual
9811	Automobile	Automobile
2368	Airplane	Airplane
9518	Truck	Truck
4599	Automobile	Ship
2648	Deer	Deer
1234	Automobile	Automobile
8059	Cat	Cat
780	Truck	Truck
6877	Ship	Airplane
131	Automobile	Automobile

9711	Deer	Deer
2102	Airplane	Airplane
2337	Horse	Horse
9285	Ship	Ship
4948	Automobile	Automobile
7863	Airplane	Airplane
5954	Ship	Ship
2398	Ship	Ship
9723	Horse	Horse
2805	Deer	Deer

```
[24]: model.save('results/model_6_2_b.h5')
```